

NoiseThresholdPrivacyMap<L01InfDistance<AbsoluteDistance<RBig>>, Approximate<PureDP>> for ZExpFamily<1>

Michael Shoemate (@Shoeboxam)

This proof resides in “**contrib**” because it has not completed the vetting process.

Proves soundness of the implementation of `NoiseThresholdPrivacyMap` for `ZExpFamily<1>` in `mod.rs` at commit `f5bb719` (outdated¹).

1 Hoare Triple

Precondition

Compiler-verified

`NoiseThresholdPrivacyMap` is parameterized as follows:

- MI, the input metric, is of type `L01InfDistance<AbsoluteDistance<RBig>>`
- MO, the output measure, is of type `Approximate<PureDP>`

Caller-verified

None

Pseudocode

```
1 # analogous to impl NoiseThresholdPrivacyMap<L01InfDistance<AbsoluteDistance<RBig>>,
   Approximate<PureDP>> for ZExpFamily<1> in Rust
2 class ZExpFamily1:
3     def noise_threshold_privacy_map(
4         self,
5         _input_metric: L01InfDistance[AbsoluteDistance[RBig]],
6         output_measure: Approximate[PureDP],
7         threshold: UBig,
8     ) -> PrivacyMap[L01InfDistance[AbsoluteDistance[RBig]], Approximate[PureDP]]:
9         #
10        noise_privacy_map = self.noise_privacy_map(
11            L1Distance.default(), output_measure[0]
12        )
13        scale = self.scale
14
15        def privacy_map(d_in: tuple[u32, RBig, RBig]):
```

¹See new changes with `git diff f5bb719..28c3d75f rust/src/measurements/noise_threshold/distribution/laplace/mod.rs`

```

16     l0, l1, li = d_in
17
18     l1_sign, l1 = l1.floor().into_parts() #
19     if l1_sign != Sign.Positive: #
20         raise f"l1 sensitivity ({l1}) must be non-negative"
21
22     li_sign, li = li.floor().into_parts() #
23     if li_sign != Sign.Positive: #
24         raise f"l-infinity sensitivity ({li}) must be non-negative"
25
26     l1 = l1.min(li * 10) #
27     li = li.min(l1) #
28
29     if l1.is_zero(): #
30         return 0.0, 0.0
31
32     if scale.is_zero(): #
33         return f64.INFINITY, 1.0
34
35     epsilon = noise_privacy_map.eval(l1) #
36
37     if li > threshold: #
38         raise f"threshold must not be smaller than {li}"
39
40     d_instability = threshold - li #
41
42     try:
43         alpha_disc = conservative_discrete_laplacian_tail_to_alpha(
44             scale, d_instability
45         )
46     except Exception:
47         alpha_disc = None
48
49     try:
50         alpha_cont = conservative_continuous_laplacian_tail_to_alpha(
51             scale,
52             d_instability,
53         )
54     except Exception:
55         alpha_cont = None
56
57     delta_single = option_min(alpha_disc, alpha_cont)
58     if delta_single is None:
59         raise "failed to compute tail bound in privacy map"
60
61     delta_joint: f64 = (1.0).inf_sub(
62         (1.0).neg_inf_sub(delta_single).neg_inf_powi(IBig.from_(10)),
63     )
64
65     # delta is only sensibly at most 1
66     return epsilon, delta_joint.min(1.0)
67
68     return PrivacyMap.new_fallible(privacy_map)

```

Postcondition

Theorem 1.1. Given a distribution `self`, returns `Err(e)` if `self` is not a valid distribution. Otherwise the output is `Ok(privacy_map)` where `privacy_map` observes the following:

Define `function(x)` as a function that updates each pair $(k_i, v_i + Z_i)$, where Z_i are iid samples from `self`, and discards pairs where $v_i + Z_i$ has smaller magnitude than $|\text{threshold}|$. The ordering of returned pairs is independent from the input ordering.

For every pair of elements x, x' in `VectorDomain<AtomDomain<IBig>>`, and for every pair (d_in, d_out) , where `d_in` has the associated type for `input_metric` and `d_out` has the associated type for `output_measure`,

if x, x' are d_{in} -close under `input_metric`, `privacy_map(d_in)` does not raise an exception, and `privacy_map(d_in) ≤ d_out`, then `function(x)`, `function(x')` are d_{out} -close under `output_measure`.

Proof. Line 9 rejects `self` if `self` does not represent a valid distribution, satisfying the error conditions of the postcondition.

We now construct the privacy map. Both the `l1` and `li` sensitivity can be floored, as neighboring integer datasets always differ in whole integer increments. This doesn't affect the check on line 23 to ensure `li` is non-negative.

There are two ways to bound the `l1` sensitivity given the three bounds:

1. The `l1` bound directly.
2. Define $x \sim x'$ as $\|x - x'\|_0 \leq 10$ and $\|x - x'\|_\infty \leq li$. Then

$$\max_{x \sim x'} \|x - x'\|_1 = \max_{x \sim x'} \sum_i^n |x_i - x'_i| \leq 10 \cdot li \quad (1)$$

Line 26 updates `l1` to the tighter of these bounds.

`li` similarly has multiple bounds:

1. The `li` bound directly.
2. Define $x \sim x'$ as $\|x - x'\|_1 \leq l1$. Then

$$\max_{x \sim x'} \|x - x'\|_\infty = \max_{x \sim x'} \max_i |x_i - x'_i| \leq \max_i [0, \dots, l1, \dots, 0] = l1 \quad (2)$$

Line 27 updates `li` to the tighter of these bounds.

Now the `l1` sensitivity is zero if any of the three bounds are zero, due to the tightening of the `l1` bound on line 26. Line 29 then checks for zero sensitivity to return a privacy loss of zero epsilon, zero delta, because all neighboring datasets have identical pairs, and the ordering is randomized.

Otherwise sensitivity is nonzero, so if the scale is zero, the privacy loss is unbounded, shown on line 32.

Now that the edge cases are handled, all sensitivities and scales are finite and strictly positive.

The mechanism that adds noise to the values in the hashmap, as described in the conditions of the postcondition of `NoiseThresholdPrivacyMap`, matches the mechanism that adds noise to a vector, as described in the conditions of the postcondition of `NoisePrivacyMap`. Therefore releasing the values incurs the privacy loss computed by `noise_privacy_map` as defined on line 35, The privacy map also ensures that `l1` is non-negative.

We now focus on computing the remaining privacy loss parameter delta. To ensure that the distance to instability remains positive, the threshold must be at least as large as the sensitivity, as checked on line 37.

Adapting the proof from [Rog23] (Theorem 7). Consider S to be the set of labels that are common between x and x' . Define event E to be any potential outcome of the mechanism for which all labels are in S (where only stable partitions are released). We then lower bound the probability of the mechanism returning an event E . In the following, c_j denotes the exact count for partition j , and Z_j is a random variable distributed according to `self`.

$$\begin{aligned} \Pr[E] &= \prod_{j \in x \setminus x'} \Pr[c_j + Z_j \leq T] \\ &\geq \prod_{j \in x \setminus x'} \Pr[\Delta_\infty + Z_j \leq T] \\ &\geq \Pr[\Delta_\infty + Z_j \leq T]^{\Delta_0} \end{aligned}$$

The probability of returning a set of stable partitions ($\Pr[E]$) is the probability of not returning any of the unstable partitions. We now solve for the choice of threshold T such that $\Pr[E] \geq 1 - \delta$.

$$\begin{aligned}\Pr[\Delta_\infty + Z_j \leq T]^{\Delta_0} &= \Pr[Z_j \leq T - \Delta_\infty]^{\Delta_0} \\ &= (1 - \Pr[Z_j > T - \Delta_\infty])^{\Delta_0}\end{aligned}$$

Let `d_instability` denote the distance to instability of $T - \Delta_\infty$.

Notice that when C is a continuous laplace random variable and D is a discrete laplace random variable,

$$\Pr[C > t] = \frac{e^{-t/s}}{2} > \frac{e^{-t/s}}{e^{1/s} + 1} = \Pr[D > t], \quad (3)$$

since $e^{1/s} > 1$. Unfortunately, the discrete laplace tail bound is not feasible to compute for large parameter choices, so we attempt to compute both tail bounds and take the minimum of the two.

By the postconditions of `conservative_discrete_laplacian_tail_to_alpha` and `conservative_continuous_laplacian_tail_to_alpha`, the tail mass is bounded above by both `alpha_disc` and `alpha_cont`. `delta_single` is the smaller of the two, as both are conservative estimates. Both bounds are computed because while the discrete bound is more accurate, it can only be computed for relatively small scales. The continuous bound is a conservative estimate for larger scales that is always computable, and converges to the discrete bound at higher scales.

The probability that a random noise sample exceeds `d_instability` is at most `delta_single`. Therefore $\delta = 1 - (1 - \text{delta_single})^{\Delta_0}$.

The privacy map returns `(epsilon, delta.min(1))`, as δ is bounded by 1. It is shown that `function(x)`, `function(x')` are `d_out`-close under `output_measure`. $\square$

References

[Rog23] Ryan Rogers. A unifying privacy analysis framework for unknown domain algorithms in differential privacy, 2023.