ApproxDPPoint::beta

Aishwarya Ramasethu (@aramasethu), Yu-Ju Ku (@yuju1998), Jordan Awan, Michael Shoemate (@Shoeboxam)

November 1, 2026

This proof resides in **"contrib"** because it has not completed the vetting process.

1 Hoare Triple

Precondition

Compiler-verified

- Receiver `self` of type `ApproxDPPoint`
- Argument `alpha` of type `RBig`

Caller-verified

`alpha` is in $[0, 1]$, and `self` was constructed by `ApproxDPPoint::build` from finite nonnegative ϵ and $\delta \in [0, 1]$.

Pseudocode

```
def beta(
    self: ApproxDPPoint,
    alpha: RBig,
) -> RBig:
    t1 = self.one_minus_delta - self.exp_eps_up * alpha
    base = max(self.one_minus_delta - alpha, RBig(0))
    t2 = self.exp_neg_eps_down * base
    return max(max(t1, t2), RBig(0))
```

Postcondition

Theorem 1.1. For all $\alpha \in [0, 1]$, `ApproxDPPoint::beta` returns a value no greater than the exact approximate-DP tradeoff curve $f_{\epsilon,\delta}(\alpha)$.

Definition 1.2. For an `ApproxDPPoint` with parameters ϵ and δ, define

$$f_{\epsilon,\delta}(\alpha) = \max(0, 1 - \delta - \exp(\epsilon)\alpha, \exp(-\epsilon)\max(1 - \delta - \alpha, 0)). \tag{1}$$

Proof. The receiver was constructed by `ApproxDPPoint::build`, which stores exact rational conversions of directed interval endpoints satisfying

$$\texttt{self.exp_eps_up} \geq \exp(\epsilon), \qquad \texttt{self.exp_neg_eps_down} \leq \exp(-\epsilon).$$

The pseudocode uses these cached constants with exact rational arithmetic.

For $\alpha \in [0, 1]$, the first affine branch is no greater than $1 - \delta - \exp(\epsilon)\alpha$ because `self.exp_eps_up` is an upper bound and α is nonnegative. The second branch is similarly no greater than its exact counterpart because `self.exp_neg_eps_down` is a lower bound. Taking the maximum with zero preserves the pointwise lower bound, so the returned value is a conservative tradeoff guarantee. □