

fn make_expr_count

Michael Shoemate (@Shoeboxam)

August 30, 2026

This document proves that the implementation of `make_expr_count` in `mod.rs` at commit `f5bb719` (outdated¹) satisfies its proof definition.

1 Hoare Triple

Precondition

Compiler-verified

- Argument `input_domain` of type `ExprDomain`
- Argument `input_metric` of type `PartitionDistance<MI>`
- Argument `expr` of type `Expr`
- Generic `MI` must implement `UnboundedMetric`
- Const generic `P` must be of type `usize`
- `(ExprDomain, PartitionDistance<MI>)` implements `MetricSpace`
- `(ExprDomain, LpDistance<P, f64>)` implements `MetricSpace`
- `Expr` implements `StableExpr<PartitionDistance<MI>, PartitionDistance<MI>`

Caller-verified

None

Pseudocode

```
1 def make_expr_count(  
2     input_domain: ExprDomain, input_metric: PartitionDistance[MI], expr: Expr  
3 ) -> Transformation:  
4     match expr: #  
5         case Agg(Count(input, include_nulls)):  
6             if include_nulls:  
7                 strategy = Strategy.Len  
8             else:  
9                 strategy = Strategy.Count  
10  
11         case Function(inputs, function=FunctionExpr.NullCount):  
12             (input,) = inputs  
13             strategy = Strategy.NullCount
```

¹See new changes with `git diff f5bb719..2bca153 rust/src/transformations/make_stable_expr/expr_count/mod.rs`

```

14
15     case Agg(NUnique(input)):
16         strategy = Strategy.NUnique
17
18     case _:
19         raise ValueError("expected count, null_count, len, or n_unique expression")
20
21     # check if input is row-by-row
22     is_row_by_row = input.make_stable(
23         input_domain.as_row_by_row(), input_metric
24     ).is_ok() #
25
26     # construct prior transformation
27     t_prior = input.make_stable(input_domain, input_metric) #
28     middle_domain, middle_metric = t_prior.output_space()
29
30     by, margin = middle_domain.context.grouping("count") #
31
32     output_domain = ExprDomain.new( #
33         column=SeriesDomain.new( #
34             middle_domain.column.name,
35             AtomDomain.default(u32),
36         ),
37         context=Context.Grouping(
38             by=by,
39             margin=Margin(
40                 max_partition_length=1,
41                 max_num_partitions=margin.max_num_partitions,
42                 max_partition_contributions=None,
43                 max_influenced_partitions=margin.max_influenced_partitions,
44                 public_info=margin.public_info,
45             ),
46         ),
47     )
48
49     match strategy: #
50     case Strategy.Len:
51         will_count_all = is_row_by_row
52     case Strategy.Count:
53         will_count_all = is_row_by_row and not middle_domain.column.nullable
54     case _:
55         will_count_all = False
56
57     public_info = (
58         margin.public_info if will_count_all else None
59     ) #
60
61     def function(e: Expr) -> Expr: #
62         match strategy:
63             case Strategy.Count:
64                 return e.count()
65             case Strategy.NullCount:
66                 return e.null_count()
67             case Strategy.Len:
68                 return e.len()
69             case Strategy.NUnique:
70                 return e.n_unique()
71
72     return t_prior >> Transformation.new( #
73         middle_domain,
74         output_domain, #
75         Function.then_expr(function),
76         middle_metric,
77         LpDistance.default(),

```

```

78     counting_query_stability_map(public_info), #
79 )

```

Postcondition

Theorem 1.1. For every setting of the input parameters (`input_domain`, `input_metric`, `expr`, `MI`, `P`) to `make_expr_count` such that the given preconditions hold, `make_expr_count` raises an error (at compile time or run time) or returns a valid transformation. A valid transformation has the following properties:

1. (Data-independent runtime errors). For every pair of members x and x' in `input_domain`, `invoke(x)` and `invoke(x')` either both return the same error or neither return an error.
2. (Appropriate output domain). For every member x in `input_domain`, `function(x)` is in `output_domain` or raises a data-independent runtime error.
3. (Stability guarantee). For every pair of members x and x' in `input_domain` and for every pair (d_in, d_out) , where `d_in` has the associated type for `input_metric` and `d_out` has the associated type for `output_metric`, if x, x' are `d_in`-close under `input_metric`, `stability_map(d_in)` does not raise an error, and `stability_map(d_in) = d_out`, then `function(x), function(x')` are `d_out`-close under `output_metric`.

Starting from line 4, `expr` is analyzed to determine the type of counting query `strategy` and input to the counting query `input`.

All preconditions for `make_stable` on line 27 are compiler-verified, therefore by the postcondition `t_prior` is a valid transformation.

To prove that the output is a valid transformation, we must first prove that the transformation on line 72 is valid.

Proof. Appropriate Output Domain Since the count transformation is not row-by-row, line 30 disallows this count transformation from being constructed in a row-by-row context, satisfying the requirements of `ExprDomain`.

By the definition of each of the allowed counting expressions in 4, the resulting output will always contain one series whose name matches the active series' name, and the data type of elements in the series is `u32`, which is consistent with 33. This series is then part of a dataframe and operations are continued to be applied with the same context, on line 32. This domain is then used on line 74 to construct the count transformation. $\square$

Proof. Stability Guarantee

To determine if `input` is row-by-row, line 24 checks if `input` can be parsed into a row-by-row transformation.

Line 49 uses domain descriptors to determine if the counting query will count all rows. This is the case if the counting query is `Len`, or if the counting query counts all non-null values in data that doesn't contain nulls.

If all rows are to be counted, then the `public_info` domain descriptor about the partitioning applies to this count query strategy, on line 59.

The function defined on line 61 applies the counting query to the input expression with the given query strategy. The stability of this function is governed by the stability map returned by `counting_query_stability_map`, as the `Len`, `Count` and `NullCount` strategies can be expressed as arbitrary predicates, and the `NUnique` strategy can be considered as a 1-stable (non-row-by-row) "unique" transformation chained with a counting query with the `Len` strategy. $\square$

Since it has been shown that both `t_prior` and the count transformation are valid transformations, then the preconditions for `make_chain_tt` are met (invoked via the right-shift operator shorthand).